

Inference cost, zero-shot vs fine-tuned (block masking, batch size 1)

